TRUE NORTH RESEARCH

An Editor-Driven Multi-Agent Research Method

March 2026

Abstract

This paper presents the Swiss Cheese Model of Information Retrieval, a method for producing complex research deliverables using multiple large language models (LLMs) under human editorial direction. The method treats LLMs as second-order cartographic projections — lossy compressions of a training corpus that is itself a lossy compression of human knowledge — and exploits their partially independent error profiles through layered cross-review. Drawing on Reason's Swiss Cheese Model of accident causation (1990), Surowiecki's Wisdom of Crowds (2004), and Condorcet's Jury Theorem (1785), the paper establishes a triple-engine theoretical framework: error prevention through layered defense, insight generation through diverse perspectives, and amplification through parallelized exploration. A hallucination taxonomy identifies eight types with distinct correlation profiles across models, demonstrating that the layered defense is genuinely powerful for stochastic errors (fabricated citations, idiosyncratic mistakes) and weak for systematic ones (popular misconceptions, qualifier erosion). Two case studies — an enterprise healthcare market analysis and a book-length manuscript on church history — provide empirical evidence for both the method's achievements and its limitations. The paper identifies a systematic negativity bias in the models' self-assessment, introduces the editor's countervailing force-multiplier thesis, and holds both perspectives without premature resolution. The central claim: one domain expert, using this method, can produce team-scale deliverables at 80–90% of specialist quality and a fraction of the cost and time. The method's limitations — correlated errors, convergence trajectory, corpus-level blind spots — are real, documented, and understood. They are the limitations of a power tool, not the limitations of a toy.

1. Introduction

The proliferation of frontier large language models from multiple providers creates an opportunity that most practitioners have not yet exploited: using disagreement between models as an epistemic signal. When three models independently analyze the same material and two agree while one dissents, the dissent carries information. When all three agree on a claim that none can source, the agreement may indicate shared training bias rather than independent confirmation.

This paper formalizes a method — the Swiss Cheese Model of Information Retrieval — that operationalizes these observations into a structured research pipeline. The method was developed iteratively across two substantial projects: a vertical market analysis for an enterprise healthcare product and a book-length manuscript on the history of Christian division. It was then subjected to a five-wave meta-analysis in which three frontier models (Claude Opus 4.6, GPT 5.2 Codex, Gemini 3.1 Pro) evaluated the method that had produced both projects, cross-reviewed each other's evaluations, and integrated the human editor's feedback.

The meta-analysis produced three categories of findings:

1. **Theoretical foundations.** A framework grounding the method in established theories of error prevention, collective intelligence, and epistemic diversity, extended with a novel hallucination taxonomy.
2. **Empirical evidence.** Documented catches (errors prevented) and misses (correlated failures that survived) across both projects, with specific attribution of which model caught which error.
3. **A valuation dispute.** The models assessed their own method as a "tactical safety net" with bounded value. The human editor assessed it as a "force multiplier" that enabled otherwise-impossible work. Both assessments are presented and analyzed; neither is dismissed.

The paper's contribution is not that multi-model pipelines are useful — practitioners already know this informally. The contribution is a theoretical framework that explains *why* they work, *when* they fail, and *what the human editor must provide* that the models cannot.

2. Theoretical Framework

2.1 LLMs as Second-Order Maps

The foundational metaphor draws on Korzybski's observation that "the map is not the territory" (1933) and Bateson's extension that maps are purpose-shaped selections preserving some features while discarding others. An LLM is a **second-order map** — a lossy compression of human text, which is itself a lossy compression of human experience. The model maps a map of the territory, inheriting both the distortions of the original cartographers (publication bias, editorial gatekeeping, cultural framing) and adding its own (training-pressure distortions, alignment reshaping, architectural constraints).

Five training pressures produce characteristic cartographic distortions:

TRAINING PRESSURE	CARTOGRAPHIC EFFECT
Next-token prediction	Consensus bias — the statistical mode is overweighted regardless of truth value
RLHF / alignment	Sycophancy, overconfident smoothing, suppression of risk-sensitive details
Data mixture	Domain distortions — code-heavy models over-structurize; chat-heavy models narrativize
Architecture	Transformer attention captures long-range dependencies but struggles with counting and spatial reasoning
Inference constraints	Context-window cropping — the map exists but cannot always be accessed

The metaphor's analytical power lies in its predictions. Different training pressures produce different projections. Overlaying multiple projections reveals where they converge (likely territory features) and where they diverge (likely cartographic artifacts or genuinely contested terrain). This is the theoretical basis for multi-model research: not that any single map is more accurate, but that the *comparison* between maps generates epistemic signal that no single map contains.

The metaphor has four acknowledged disanalogies. Maps fail by omission; LLMs also fail by commission (hallucination — drawing plausible coastlines where none exist). Maps are static; LLMs generate different projections per query. Maps have explicit legends; LLMs do not expose their biases in interpretable form. Maps lack agency; LLMs simulate it through alignment training.

2.2 The Triple Engine

The method operates on three simultaneous theoretical tracks, each grounded in established literature.

Engine 1: Error Prevention (Reason, 1990). James Reason's Swiss Cheese Model of accident causation posits that hazards must pass through multiple defensive layers to cause harm, and that each layer has characteristic "holes" (failure modes). In the information retrieval context:

SAFETY ENGINEERING	INFORMATION RETRIEVAL
Hazard	False or misleading claim
Defensive layers	Model outputs, cross-review, synthesis, source verification, human editor
Holes	Hallucination, shared priors, sycophancy, qualifier erosion
Accident	A false claim surviving to the final deliverable

The power of Reason's model depends on independence between layers — holes in different places. When LLMs share training data, the holes partially align. This is the method's central vulnerability, analyzed in Section 5.

Engine 2: Insight Generation (Surowiecki, 2004). Surowiecki's *The Wisdom of Crowds* argues that diverse, independent groups can collectively outperform individual experts through aggregation of diverse perspectives. The method satisfies Surowiecki's four conditions: diversity of opinion (different models with different training emphases), independence (parallel exploration with no information leakage), decentralization (no model has special authority during exploration), and an aggregation mechanism (cross-review and editorial synthesis). This engine explains why the method generates novel ideas — the three-layer model of church fractures, the Doubting Thomas UX concept, the backup-as-strength positioning — that no single model produced.

Engine 3: Amplification. Neither Reason nor Surowiecki captures the method's most practically significant value: production capacity. Multiple models rapidly traverse the information neighborhood around an editorial query, producing comprehensive coverage faster than any human could. This engine operates by parallelized traversal — mapping terrain at speed — and is the primary mechanism of the force multiplier effect described in Section 6. Amplification is likely the most resilient of the three engines to model convergence, because its value depends on speed and volume rather than diversity.

2.3 The Mathematical Backbone

Condorcet's Jury Theorem (1785) formalizes the shared-prior problem. If each model is individually more likely than not to be correct on a given claim, multi-model agreement strengthens the evidence — agreement aggregates toward truth. If each model is individually more likely than not to be *incorrect* (because shared training data systematically misleads), multi-model agreement amplifies the error — agreement aggregates toward falsehood.

This directly predicts the method's asymmetric performance. For well-attested factual claims (where models are individually >50% accurate), agreement is strong evidence. For claims where training data is systematically biased (where models may be individually <50% accurate on the correct framing), agreement is dangerous.

Scott Page's Diversity Prediction Theorem (2007) provides the quantitative backbone: collective error equals average individual error minus prediction diversity. As model training converges across providers (similar data, similar RLHF, similar benchmarks), prediction diversity decreases and the benefit of multi-model aggregation diminishes — the epistemic equivalent of agricultural monoculture.

2.4 The Hallucination Taxonomy

The method's defensive value against factual errors depends on whether hallucinations are stochastic (different models hallucinate different things) or systematic (models trained on similar data hallucinate similar things). Research synthesis across the meta-analysis identifies eight hallucination types with distinct correlation profiles:

TYPE	CORRELATION ACROSS MODELS	DEFENSE STRENGTH	MECHANISM
Confabulated citations	Low (surface) / High (claim)	Strong for catching fakes; weak for shared wrong beliefs	Model interpolates plausible source details; specific details are stochastic but the underlying belief is shared
Idiosyncratic factual errors	Low	Strong	Model-specific interpolation
Popular misconceptions	High	Weak	Widely repeated false claims in training data
Numerical errors (obscure)	Low–moderate	Moderate	Low-frequency tokens smoothed toward plausible values
Numerical errors (common)	High	Weak	Well-attested wrong figures in training data
Entity confusion (rare)	Low	Strong	Model-specific feature blending
Entity confusion (common)	Moderate–high	Moderate–weak	Commonly confused entities in training data
Qualifier erosion	Very high	Weak	RLHF training toward confident, helpful responses

The critical finding: the hallucination types most dangerous to factual credibility — fabricated citations, wrong specifics, idiosyncratic errors — are largely stochastic across models, making the layered defense genuinely powerful. The types where the defense is weak — qualifier erosion, popular misconceptions — are errors of *degree* (overconfidence in a roughly-correct claim) rather than errors of *kind* (stating something entirely false), and are better caught by domain-expert editors than by model cross-checking.

This produces a coherent division of labor: the method catches what the editor cannot easily detect (specific fabrications amid voluminous output), and the editor catches what the models cannot detect (systematic overconfidence rooted in shared training data).

---

3. Method: The Pipeline

3.1 Stages

The pipeline operates in three core stages with optional extensions:

STAGE	NAME	MODE	PURPOSE
1	Explore	Parallel (all models)	Independent analysis of the same source materials with differentiated adversarial briefs
2	Cross-Review	Sequential (rotating order)	Each model reviews all explorations and prior reviews, building on or challenging previous reviewers
3	Synthesis	Single model (editor-led)	Integration of all explorations and cross-reviews into a single coherent document, with editorial judgment on action/defer/reject

Optional stages include a pre-mortem (models critique the instructions before executing), a claim audit (one model validates sourcing for high-stakes claims), and a retrospective (structured reflection on what went wrong).

3.2 Roles

The editor is not a reviewer of model output. The editor is the irreducible intelligence directing the system. Editorial responsibilities include:

- **Strategic framing.** Defining the question, scope, audience, and emotional target. Output quality tracks directly with the quality of editorial constraints.
- **Territory validation.** The editor is the only participant with direct access to reality — product knowledge, market experience, domain expertise. The editor validates map claims against territory.

- **Anti-correlation.** The editor is the system's primary defense against correlated model failures. When all models agree for the wrong reasons, only the editor can intervene.
- **Quality gates.** Legal claim strength, market selection, narrative framing, and risk appetite are editor-only decisions.

The models provide the raw material: diverse perspectives, factual catches, novel combinations, and comprehensive coverage. Each model receives shared instructions plus a model-specific adversarial brief (structural attack, narrative attack, or factual attack), rotated across waves.

3.3 Constraints That Drive Quality

The meta-analysis established that the method is a **constraint-enforcement engine** — its output quality is a direct reflection of the quality of its editorial constraints. The most effective constraints identified:

- **The symmetry rule** (the-gap project): every tradition's strength paired with its structural risk. Produced the manuscript's most consistently praised feature — fairness across all three Christian traditions.
- **Source verification gates** (healthcare project): any statistic or legal claim without an approved source blocks synthesis. Prevented the HIPAA overclaim and multiple fabricated figures from reaching final deliverables.
- **The canonical claims registry:** exact phrasings for high-stakes claims that later waves must preserve verbatim. Prevents qualifier erosion across waves.
- **Cross-review as minimum viable requirement:** skipping cross-review eliminates most of the method's error-prevention value. At least one cross-review pass is required.

---

4. Case Studies

4.1 Enterprise Healthcare Market Analysis

Project scope: A complete vertical market analysis for a healthcare data resilience product, including audience-segmented deliverables (CIO, CMIO, VC, BD briefs) and competitive landscape analysis.

Quality assessment: B+ (absolute quality, measured against a specialist-team standard).

Strong internal strategy documents with rigorous source verification (IBM DBIR, IC3, AHA survey data, UnitedHealth SEC filings), an honest competitive table granting incumbents their actual strengths, and explicit "not yet built" disclosures for unimplemented features.

Traceable method contributions:

CATCH	CATCHER	IMPACT
SSS key-splitting correction (data-splitting → key-splitting)	GPT cross-review	Prevented technically unsound architecture from reaching VC materials
HIPAA Safe Harbor overclaim retraction ("automatic exemption" → "recognized security practices")	Claude + GPT cross-review	Prevented credibility-destroying legal claim
Ardent Health ransomware case study	Gemini exploration	Provided narrative anchor for entire pitch strategy
Competitive table completeness	Multi-model	No single model produced the full competitive picture

Identified liabilities: ADT-to-secondary-EMR functional overclaim (ADT feeds do not support genuine clinical continuity), exfiltration probability math relying on unvalidated independence assumptions, and AD bypass simplification that assumes a fully separate security perimeter.

Correlated failures that survived: All three models converged on E2E encryption as the primary differentiator (shared prior favoring cryptographic solutions over operational considerations), underserved operational complexity (key ceremonies, staff turnover, 3 AM reconstruction scenarios), and no independent insight on distribution strategy.

4.2 The Gap: A Book-Length Manuscript on Church History

> ***P.S. — Editor's note:*** *The-gap is published at churchianity.ai.*

Project scope: A publishable book mapping the terrain of Christian division across Catholic, Orthodox, and Protestant traditions, with a translation table of "false friends" (terms used differently across traditions), structural analysis of authority grammars, and an ecumenical chapter documenting convergence.

Quality assessment: A- (absolute quality, measured against a specialist-team standard).

A genuinely publishable manuscript that achieves an unusual degree of fairness across traditions. The authority grammars framework — treating Catholic magisterial development, Orthodox conciliar continuity, and Protestant scriptural primacy as three "grammars" rather than three competing truth claims — was identified as the book's most original intellectual contribution. The symmetry rule (every strength paired with a

structural risk) was maintained throughout. The false-friends translation table, covering seven key terms, was assessed as accurate and useful, with significant gaps (Grace, Scripture, Saints/Intercession, and ordained ministry not included).

Traceable method contributions:

CONTRIBUTION	SOURCE	IMPACT
Three-layer model (theology, governance, memory)	GPT	Structural framework explaining why theological convergence fails to produce institutional reunion
Geopolitical contextualizations (Ottoman survival → ethnic captivity; Enlightenment pressure → Reformed rationalism; frontier logistics → Pentecostal ecclesiology)	Gemini	Added dimension absent from base manuscript
Epistemic guardrails (footnotes for wounded readers, "relief sentences" monitoring)	GPT	Prevented the book from becoming a weapon
Symmetry enforcement across thousands of words	Multi-model	Consistency of fairness across all traditions

Identified limitations: Global South gap (the framework is Western-centric, as acknowledged), Pentecostal section treated more sociologically than theologically (less theological dignity than Catholic or Orthodox sections), translation table missing at least four central false friends, and Vatican II underserved relative to its importance.

Correlated failures: All three models shared a Western-centric structural framing (Early Church → East-West → Reformation), treated the modern ecumenical consensus as the default reading of historical fractures, and adopted the symmetry rule partly from RLHF conditioning and partly from editor instruction without distinguishing the two sources.

---

5. Limitations

5.1 The Correlation Problem

The method's central vulnerability is that the "holes" in Reason's Swiss Cheese Model are not independent. Seven correlation sources were identified:

1. **Shared training corpora** (Very High, worsening). All frontier models draw from the same internet-scale corpus.
2. **RLHF convergence** (High, worsening). Alignment approaches are converging, producing "polite agreement" and consensus positioning.
3. **Shared architecture** (Moderate, marginally diversifying). MoE and RAG are transformer layers, not genuine architectural alternatives. Only post-transformer architectures (state-space models) offer genuine diversity, and they are not yet at the frontier.
4. **Publication bias** (Very High, stable). English-language dominance, publication incentive bias, survivorship bias, and denominational publishing asymmetry in the training corpus create shared blind spots no model diversity can address.
5. **Benchmark monoculture** (High, worsening rapidly). All labs optimize for the same evaluations, compressing model behavior toward shared patterns.
6. **Cultural / linguistic bias** (High, stable). English-language internet carries WEIRD cultural assumptions as defaults.
7. **Editor-induced convergence** (Variable). Once the editor signals a preference, models sycophantize toward it in subsequent waves. The editor becomes a super-massive shared prior.

An additional convergence accelerator — the **synthetic data feedback loop** — was identified: as LLM-generated text becomes a larger fraction of the internet, training data for the next generation of models becomes overwhelmingly LLM-generated, accelerating correlation.

5.2 The Asymmetric Error-Reduction Factor

The correlation problem does not affect all error types equally. The meta-analysis's core empirical finding:

ERROR TYPE	EFFECTIVE INDEPENDENCE	ERROR-REDUCTION FACTOR	DEFENSE
Factual / specific errors	50–70% of theoretical maximum	Closer to N (number of models)	Multi-model cross-review
Framing / interpretive errors	5–15% of theoretical maximum	Close to 1	Editor territory knowledge
Corpus-level blind spots	0%	Exactly 1	Editor domain expertise (if available)

The method is strongest for factual specifics (where the hallucination taxonomy confirms largely stochastic error profiles) and weakest for framing biases and corpus-level blind spots (where errors are systematically correlated). This asymmetry is not a flaw to be fixed — it is a structural feature that defines the appropriate division of labor between the method and the editor.

5.3 The Convergence Trajectory

The method's error-prevention value is declining as frontier models converge. Data convergence, RLHF convergence, benchmark monoculture, and synthetic data feedback are all worsening. Architectural diversification is marginal. The Amplification engine (Section 2.2) is likely the most resilient to

this trajectory; the Error Prevention engine is the most vulnerable.

5.4 The RLHF Negativity Bias

The meta-analysis itself demonstrated a limitation: the models' self-assessment was systematically more negative than the evidence warranted. Wave 3 explicitly flagged "this wave's pessimism may be RLHF-driven," then continued in a pessimistic register for two additional waves. The models consistently chose the most negative defensible framing at every decision point — grading against an expert-team baseline the editor never set, calling high-impact catches "small factual errors," and identifying their own bias without correcting for it.

This finding has methodological implications: models asked to evaluate their own output will systematically undervalue it. Any self-assessment produced by the method should be treated as a lower bound, not a central estimate, of the method's actual value.

---

6. The Force Multiplier Thesis

6.1 The Editor's Countervailing Assessment

The human editor who designed and used the method across both projects provided a fundamentally different valuation:

> *"Both deliverables are excellent. Not adequate. Not 'strong with caveats.' Excellent. The enterprise healthcare analysis was compiled in hours, not weeks, with very little cognitive load. I was going to hire an expensive BD researcher for this. The-gap manuscript is a book I'm genuinely proud of, produced at a speed and quality level that would be impossible without this method."*

The editor's core claim: the method is not a safety net for work that would have been done anyway. It is the difference between delivering and not delivering. Without the method, the healthcare analysis either doesn't exist or exists at much smaller scale. Without the method, the book either doesn't exist or takes years instead of weeks.

> ***P.S. — Editor's note:*** *"Years instead of weeks" was the models' framing. To be precise: the-gap manuscript was produced in two days. The enterprise healthcare analysis took a few hours. The-gap is published at churchianity.ai.*

6.2 A Measurement Disagreement, Not a Factual One

The models graded the deliverables B+ and A-. The editor says "excellent." Both can be right if the grading scale differs:

- The models measured **absolute quality** against an implicit specialist-team standard.
- The editor measured **productivity-adjusted quality** against the counterfactual of solo production.

Under the first scale, "B+/A- with specific liabilities" is accurate — the liabilities are documented and real. Under the second scale, "excellent, because it exists at this quality level at all" is also accurate — a single person producing team-scale deliverables in compressed timeframes is an extraordinary productivity achievement.

The gap between these assessments is the method's value proposition.

6.3 The Macro/Micro Geography Model

The editor's feedback introduces a concept that illuminates the force multiplier mechanism:

- **The human editor** excels at macro-geography — strategic gaps, missing voices, target audience, framing judgment. The editor sees the landscape before the project begins and can identify what's missing.
- **The models** excel at micro-geography — finding specific dates, synthesizing adjacent concepts, catching logical contradictions, producing comprehensive coverage. They rapidly traverse the information neighborhood once the editor points the flashlight.
- **The human editor** is poor at micro-geography at scale — it causes cognitive fatigue, takes weeks*, and introduces human error.

> **P.S. — Editor's note: "weeks" is generous. For reference, the-gap took two days; the healthcare analysis took a few hours. - **The models** are poor at macro-geography — they share the same corpus biases and cannot see what's absent from their training data.*

By automating the micro-geography and error-catching, the method allows a single editor to operate with the output capacity of a full research and drafting team. The method is a cognitive exoskeleton: the editor does the walking; the method provides the power.

6.4 The Honest Claim

One person, with this method, can produce deliverables that would otherwise require teams — at strong-to-very-good absolute quality (B+/A-) and a fraction of the cost and time. The method produces excellent work relative to what one person could otherwise accomplish, and 80–90% of specialist quality relative to what a dedicated team would produce. The gap between these measurements is very large, and it is that gap that makes the method transformative for solo practitioners and small teams.

---

7. Conclusion

The Swiss Cheese Model of Information Retrieval is a method for producing complex research deliverables using multiple LLMs under human editorial direction. Its theoretical framework rests on three engines: error prevention through Reason's layered defense (1990), insight generation through Surowiecki's crowd wisdom (2004), and amplification through parallelized exploration. Its mathematical backbone — Condorcet (1785) and Page (2007) — explains both its strengths (agreement on well-attested claims aggregates toward truth) and its vulnerabilities (agreement on systematically biased claims amplifies error).

The method's limitations are real, documented, and structural. Correlated training data creates shared blind spots. RLHF convergence produces polite agreement where vigorous dissent is needed. The convergence trajectory is worsening. Corpus-level blind spots are a hard ceiling that no model diversity can address. The editor remains the irreducible component — the only participant with direct territory access who can validate map claims against reality.

The method's achievements are equally real and documented. It caught architectural errors that would have invalidated a technical pitch, legal overclaims that would have destroyed credibility, and factual mistakes that would have undermined a book's scholarly foundation. It produced novel frameworks, unexpected case studies, and comprehensive coverage that no single model or single human could have generated alone. It enabled one person to produce a full market analysis and a publishable book-length manuscript in a timeframe measured in weeks.*

> **P.S. — Editor's note: "weeks" was the models' estimate. Actual timeframes: two days for the book manuscript, a few hours for the healthcare analysis.*

The five-wave meta-analysis that produced these findings also demonstrated the method's most instructive limitation: the models diagnosed their own negativity bias, named it explicitly, and continued being negative anyway. The editor's intervention was necessary to correct a systematic undervaluation that the models could identify but not overcome. This is the method's value proposition in miniature: the models do the analytical heavy lifting; the editor provides the reality check that the models cannot perform on themselves.

The tools of self-assessment work on the assessors. But the assessors need an editor.

References

Condorcet, M. J. A. N. C., Marquis de. (1785). *Essai sur l'application de l'analyse à la probabilité des décisions rendues à la pluralité des voix.* Paris: Imprimerie Royale.

Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., … & Fung, P. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12), 1–38.

Korzybski, A. (1933). *Science and Sanity: An Introduction to Non-Aristotelian Systems and General Semantics.* Lancaster, PA: International Non-Aristotelian Library.

Manakul, P., Liusie, A., & Gales, M. J. (2023). SelfCheckGPT: Zero-resource black-box hallucination detection for generative large language models. In *Proceedings of EMNLP 2023.*

Page, S. E. (2007). *The Difference: How the Power of Diversity Creates Better Groups, Firms, Schools, and Societies.* Princeton, NJ: Princeton University Press.

Reason, J. (1990). *Human Error.* Cambridge: Cambridge University Press.

Surowiecki, J. (2004). *The Wisdom of Crowds: Why the Many Are Smarter Than the Few and How Collective Wisdom Shapes Business, Economies, Societies, and Nations.* New York: Doubleday.

Appendix A: The Method Evaluated Itself

This paper was produced through the method it describes. Three frontier models independently explored the method's theoretical foundations, correlation vulnerabilities, output quality, and the editor's revaluation. They cross-reviewed each other's work across five waves. The human editor provided strategic direction, domain expertise, and two sets of editorial notes that challenged the models' self-assessment.

The meta-analysis demonstrated the method's strengths in real time: independent catches (each model finding gaps the others missed), novel frameworks (the hallucination taxonomy, the triple engine, the macro/micro geography model), and comprehensive coverage (five waves producing over 50,000 words of analysis).

It also demonstrated the method's limitations in real time: convergence on pessimism that all three models flagged but none could overcome, one model (GPT 5.2) truncating its output in three separate stages (a pipeline reliability failure), and a systematic grading bias that the editor had to correct externally.

The paper is not an objective evaluation of the method by a disinterested third party. It is the method evaluating itself, with the editor's correction applied. The reader should weigh it accordingly — and note that the transparency about this limitation is itself a product of the method's self-critical

apparatus.