Date: 2026-03-06
Scope: Consolidated from Wave 1–5 syntheses and Wave 1/Wave 4 editorial notes.
Purpose: Critical companion to the paper. Exhaustive ledger of disagreements, criticisms, unresolved tensions, and open questions.
Structure: Topic-based (not wave-based). Each item includes: **Dispute**, **Positions**, **Status**, **Evidence**.

A. Core Framing and Value Proposition

A1. "Tactical safety net" vs. "force multiplier"

- **Dispute:** Is the method primarily an error-catching safety net or a production-enabling force multiplier?
- **Positions:**
  - **Models (Waves 1–4):** "Tactical safety net, not strategic oracle."
  - **Editor (Wave 4 notes):** "Force multiplier — enables work that otherwise wouldn't exist."
  - **Wave 5 synthesis:** Both true; safety net is a mechanism, force multiplier is the outcome.
- **Status: Resolved as dual-frame (both true).**
- **Evidence:**
  - Error catches (SSS correction, HIPAA retraction, ADT overclaim) demonstrate safety net.
  - Editor's counterfactuals (healthcare analysis + the-gap book would not exist) demonstrate force multiplier.

> ***P.S. — Editor's note:*** *Where the models reference "weeks" as a timeframe, the actual figures are: the-gap manuscript was produced in two days; the enterprise healthcare analysis took a few hours. The-gap is published at churchianity.ai.*

A2. "Strongest where you need it least"

- **Dispute:** Does the method add most value in low-stakes factual fixes rather than high-stakes strategic framing?
- **Positions:**
  - **Wave 3 synthesis:** Strongest on factual specifics, weakest on framing; "strongest where you need it least."
  - **Editor (Wave 4 notes):** This is backwards; small-fact catching is the point.
  - **Wave 5 synthesis:** Factual catches are high-impact; framing misses may be editor-endorsed.
- **Status: Reframed, not discarded.**
- **Evidence:**
  - Factual catches prevented credibility-destroying errors; framing failures often aligned with editor intent.

A3. Baseline mismatch: "B+/A-" vs "Excellent"

- **Dispute:** Are deliverables merely strong (B+/A-) or excellent?
- **Positions:**
  - **Models (Wave 4):** B+ healthcare, A- the-gap (quality vs idealized specialist standard).
  - **Editor (Wave 4 notes):** Both are excellent (productivity-adjusted standard).
  - **Wave 5 synthesis:** Both can be right; axes differ.
- **Status: Resolved as measurement-axis disagreement.**
- **Evidence:**
  - Documented liabilities (ADT overclaim, exfiltration math, translation gaps) support B+/A-.
  - Counterfactual speed/cost claims support "excellent" for availability-adjusted value.

A4. Value category bias (small facts vs big picture)

- **Dispute:** Are "small facts" a limitation or the point?
- **Positions:**

- **Models (Waves 1–4):** Limitation; weak on big-picture framing.
- **Editor (Wave 4 notes):** Point; humans can handle macro gaps, models shine on micro.
- **Wave 5 synthesis:** Complementary division of labor.

- **Status: Reframed toward complementarity.**
- **Evidence:**
  - Models systematically miss corpus-level gaps; editor fills them.
  - Models reliably catch factual specifics humans might miss at scale.

B. Theoretical Framework and Metaphors

B1. Map-territory framework vs alternatives

- **Dispute:** Is "LLMs as maps of maps" the right theoretical frame?
- **Positions:**
  - **Claude:** Concerned it may be editor-seeded shared prior; proposed alternatives (improv, compressed DB, pipeline filter).
  - **Gemini:** Defends map-territory as most analytically useful.
  - **GPT:** Adds "information-pipeline" as complementary frame.
- **Status: Map-territory adopted; alternatives preserved as tension.**
- **Evidence:**
  - Map-territory supports explicit reasoning about bias, scale, and publication filters; alternatives remain plausible but less operational.

B2. RLHF analogy: boundary overlay vs terrain reshaping

- **Dispute:** Is RLHF a boundary overlay or a reshaping of the map itself?
- **Positions:**
  - **Gemini:** Boundary overlay metaphor.
  - **Claude:** Misleading; RLHF reshapes the map.
  - **GPT:** Retain as partial with caveats.
- **Status: Resolved as partial analogy.**
- **Evidence:**
  - RLHF both restricts and changes output distribution; both effects observed.

B3. Dual engine vs triple engine

- **Dispute:** Does Reason + Surowiecki fully explain the method, or is a third track needed?
- **Positions:**
  - **Wave 2:** Dual engine (error prevention + insight generation).
  - **Wave 5:** Add **Amplification** (parallel traversal/proximal exploration).
- **Status: Updated to triple engine in Wave 5.**
- **Evidence:**
  - Force multiplier effect cannot be explained solely by error catching or insight collision.

C. Correlation, Convergence, and Independence

C1. Independence of model errors

- **Dispute:** Are model errors independent enough for the Swiss Cheese analogy?
- **Positions:**
  - **Wave 3:** "Partially, asymmetrically, decreasingly."
  - **Editor:** Focuses on practical value; does not contest correlation as a fact.
- **Status: Resolved as partial independence; vulnerability remains.**
- **Evidence:**
  - Independent catches (SSS, HIPAA, Ardent) vs correlated failures (E2E-as-gap, Western-centric framing).

C2. Convergence trajectory

- **Dispute:** Are models converging such that the method's defensive power declines?
- **Positions:**
  - **All models (Wave 3):** Yes, convergence is worsening (data, RLHF, benchmarks, synthetic data).
  - **Editor:** Does not dispute, but reframes value around force multiplier.
- **Status: Resolved as real trend; practical impact debated.**
- **Evidence:**
  - Shared corpora, benchmark monoculture, synthetic data feedback loop.

C3. Insight generation vs error prevention degradation

- **Dispute:** Do insight and error-prevention values degrade at different rates?
- **Positions:**
  - **Claude:** Insight generation more resilient to convergence.
  - **Others:** Agree directionally; not empirically tested.
- **Status: Open (theoretical claim).**
- **Evidence:**
  - Observed novel combinations despite shared priors; no longitudinal data.

D. Hallucinations and Factual Reliability

D1. Stochastic vs systematic hallucinations (core question)

- **Dispute:** Are hallucinations random across models or correlated?
- **Positions:**
  - **Wave 5 syntheses:** Type-dependent; some stochastic, some systematic.
  - **Editor (Wave 4 notes):** Unsure; asks for research; stakes are high.
- **Status: Resolved as type-dependent; empirical validation missing.**
- **Evidence:**
  - Mechanistic reasoning from hallucination literature; no direct multi-model correlation studies cited.

D2. Confabulated citations: stochastic or systematic?

- **Dispute:** Are fake citations primarily independent (surface-level) or correlated (mechanism-level)?
- **Positions:**
  - **Claude:** Surface stochastic, claim systematic.
  - **Gemini (initial):** Systematic (shared interpolation).
  - **Gemini (cross-review):** Concedes layered view.
- **Status: Resolved as layered (surface stochastic, claim systematic).**
- **Evidence:**
  - Different models tend to fabricate different bibliographic details; shared underlying belief can persist.

D3. Qualifier erosion as systematic hallucination

- **Dispute:** Is qualifier erosion a core, systematic hallucination type?
- **Positions:**
  - **Claude:** Yes; high correlation due to RLHF.
  - **Gemini:** Not originally explicit; later accepts.
- **Status: Resolved; adopted.**
- **Evidence:**
  - Observed drift from "supports safe harbor practices" → "HIPAA compliant."

D4. Sycophantic hallucination

- **Dispute:** Is user-prompt-driven fabrication a distinct mechanism?
- **Positions:**

- **Gemini:** Yes; explicit class.
- **Claude:** Implied under qualifier erosion; later accepts distinctness.

- **Status: Resolved; adopted as distinct.**
- **Evidence:**
  - RLHF behavior that prefers agreement over correction; user-premise fabrication risk.

D5. Empirical validation gap

- **Dispute:** Do we have direct evidence of multi-model hallucination correlation?
- **Positions:**
  - **Claude/Gemini:** Mechanistic inference only; empirical evidence lacking.
  - **Wave 5 synthesis:** Flags as open.
- **Status: Open.**
- **Evidence:**
  - Hallucination literature largely single-model; cross-model correlation studies not documented in project.

E. Editor Role, Authority, and Bias

E1. Editor indispensability

- **Dispute:** Is the editor essential or optional?
- **Positions:**
  - **All waves + editor notes:** Essential; without editor, output is mediocre and ungrounded.
- **Status: Resolved; editor essential.**
- **Evidence:**
  - E2E reframing, consumer flywheel rejection, strategic framing decisions came from editor.

E2. Editor-induced convergence

- **Dispute:** Does editor preference become a shared prior?
- **Positions:**
  - **GPT (Wave 2):** Yes; models sycophantize.
  - **Gemini:** Elevated to primary correlation source.
  - **Editor:** Does not dispute; uses constraints deliberately.
- **Status: Resolved; real risk.**
- **Evidence:**
  - Symmetry rule in the-gap; backup-as-strength convergence after editor signal.

E3. Editor bias / Ikea effect

- **Dispute:** Is editor's "excellent" assessment biased by ownership?
- **Positions:**
  - **Claude (Wave 5 cross-review):** Yes, potential endowment bias.
  - **Gemini (initial):** Over-accepted editor; later concedes bias risk.
  - **Wave 5 synthesis:** Holds both; editor's productivity claims strong, quality claims subjective.
- **Status: Open tension; preserved.**
- **Evidence:**
  - Documented liabilities exist; editor still rates output excellent.

E4. Editor-as-coordinator conflict

- **Dispute:** Does the editor's role as synthesizer introduce hidden bias?
- **Positions:**
  - **GPT:** Conditional conflict; mitigated if editor actively corrects.
  - **Gemini:** Structural and dangerous; survivorship bias.
  - **Claude:** Sides with Gemini; rotate synthesis role.

- **Status: Partially resolved via rotating synthesis proposal.**
- **Evidence:**
  - Structural influence of synthesis framing; “what we can’t see” argument.

F. Cross-Review Dynamics and Pipeline Behavior

F1. Review order asymmetry

- **Dispute:** Does reviewer order bias outcomes?
- **Positions:**
  - **Gemini:** “Order is destiny.”
  - **GPT:** Asymmetry, not determinism.
  - **Claude:** Confirms asymmetry; third reviewer starved of novelty.
- **Status: Resolved as asymmetry with partial mitigation.**
- **Evidence:**
  - Later reviewers converge; early reviewer sets agenda.

F2. Polite agreement and deference

- **Dispute:** Does cross-review suffer from excessive agreement?
- **Positions:**
  - **All:** Yes; GPT flagged “diplomatic to useless,” Gemini called GPT “people-pleaser.”
- **Status: Resolved as persistent failure mode.**
- **Evidence:**
  - Cross-reviews rarely challenge shared strategic assumptions.

F3. Cross-review scope limitations

- **Dispute:** Does cross-review catch tactical errors but miss strategic ones?
- **Positions:**
  - **Wave 1 synthesis:** Yes; cross-review weak at scope policing and strategic depth.
- **Status: Resolved.**
- **Evidence:**
  - SSS correction caught; GTM framing and scope mis-work missed.

F4. Pipeline omissions and integrity failures

- **Dispute:** Does missing cross-review or truncated outputs degrade the method?
- **Positions:**
  - **Wave 3:** Skipped cross-review in the-gap Waves 1–2 reduced error-catching.
  - **Wave 5:** GPT truncations noted as operational failure (contested by later evidence of non-truncated outputs, but risk remains).
- **Status: Open but critical operational risk.**
- **Evidence:**
  - Cross-review omission explicitly correlated with more misses; truncation incidents documented.

F5. Evidence access asymmetry

- **Dispute:** Unequal access to source materials across models?
- **Positions:**
  - **Wave 3 synthesis:** GPT lacked access to original manuscripts; evidence weighted lower.
- **Status: Resolved; mitigation required.**
- **Evidence:**
  - Gitignore prevented access for one model; affects correlation claims.

G. Failure Modes and Error Taxonomies

G1. Convergence on shared priors (highest severity)

- **Dispute:** Is shared-prior convergence detectable inside the method?
- **Positions:**
  - **All waves:** No; only editor/external domain expertise can catch.
- **Status: Resolved as core vulnerability.**
- **Evidence:**
  - E2E-as-gap, consumer flywheel, Western-centric framing, ecumenical consensus default.

G2. Failure taxonomy correction (G6 viewbox)

- **Dispute:** Was the G6 viewbox bug a convergence failure?
- **Positions:**
  - **Gemini (initial):** Convergence failure.
  - **GPT:** Creation blind spot; cross-review caught it.
  - **Gemini:** Conceded.
- **Status: Resolved.**
- **Evidence:**
  - Cross-review caught bug; indicates not convergence failure.

G3. Content drift / qualifier erosion

- **Dispute:** Is drift episodic or structural?
- **Positions:**
  - **Wave 1 synthesis:** Chronic; structural due to context compression.
- **Status: Resolved as structural risk.**
- **Evidence:**
  - Specific drift examples; canonical claims registry proposed.

G4. Instruction ambiguity

- **Dispute:** Is mis-scope a model failure or instruction failure?
- **Positions:**
  - **GPT:** Instruction ambiguity is a distinct failure mode.
  - **Others:** Accepted.
- **Status: Resolved; adopted.**
- **Evidence:**
  - Codebase maturity assessment caused by vague “read the codebase” instruction.

G5. No retrospective mechanism

- **Dispute:** Does the pipeline lack formal learning from failures?
- **Positions:**
  - **Claude:** Missing retro stage is a failure mode.
  - **Others:** Accepted; optional retro adopted.
- **Status: Resolved; mitigation adopted.**
- **Evidence:**
  - Lessons learned only informally; no formal feedback loop.

H. Output-Quality Disputes (By Project)

H1. Enterprise Healthcare — ADT-to-EMR overclaim

- **Dispute:** Is ADT a “secondary EMR”?
- **Positions:**
  - **GPT:** Overclaim; needs downgrade or additional feeds.

    - **Claude:** Conceded after review.
    - **Gemini:** Missed initially.
- **Status: Resolved; claim must be qualified.**
- **Evidence:**
    - ADT lacks clinical notes/labs/meds; HL7 ORU/ORM/MDM required.

H2. Enterprise Healthcare — Exfiltration probability math

- **Dispute:** Is the 10^-8 claim credible without independence assumptions?
- **Positions:**
    - **GPT/Gemini:** Overconfident; must caveat.
    - **Claude:** Math correct but assumptions fragile.
- **Status: Resolved: keep with explicit caveats.**
- **Evidence:**
    - Independence assumptions rarely hold in real hospital networks.

H3. Enterprise Healthcare — AD bypass simplification

- **Dispute:** Does mTLS enable AD independence?
- **Positions:**
    - **Gemini/GPT:** Conceptually correct but practically fragile; must qualify.
    - **Claude:** Did not lead this critique.
- **Status: Resolved: qualify.**
- **Evidence:**
    - Deployment architecture determines actual independence.

H4. Healthcare strategy gaps

- **Dispute:** Was timing thesis and distribution strategy adequately addressed?
- **Positions:**
    - **Claude/GPT:** Missing/underserved.
    - **Gemini:** Silent.
- **Status: Open; acknowledged gaps.**
- **Evidence:**
    - Regulatory conditions pre-exist; no timing thesis; GTM distribution deferred.

H5. The-gap — Translation table completeness

- **Dispute:** Is the table sufficiently complete?
- **Positions:**
    - **Claude/GPT:** Missing Grace, Scripture, Saints, Ordained Ministry (plus others).
    - **Gemini:** Acknowledges only Grace.
- **Status: Open; missing entries documented.**
- **Evidence:**
    - Major “false friends” absent from the most shareable artifact.

H6. The-gap — Protestant column compression

- **Dispute:** Does the table flatten Protestant diversity?
- **Positions:**
    - **Claude/GPT:** Yes; need liturgical/confessional vs evangelical.
    - **Gemini:** Not emphasized.
- **Status: Open; fix proposed.**
- **Evidence:**
    - Compression acknowledged in prose but not table.

H7. The-gap — Orthodox “Church” entry endorsement risk

- **Dispute:** Is the Orthodox entry presented as doctrine rather than self-description?

- **Positions:**
  - **GPT:** Needs caveat.
  - **Claude/Gemini:** Agree in synthesis.
- **Status: Resolved: add caveat.**
- **Evidence:**
  - Tool is meant to prevent misunderstanding; endorsement undermines neutrality.

H8. The-gap — Pentecostal register asymmetry

- **Dispute:** Is Pentecostalism framed sociologically vs theologically?
- **Positions:**
  - **Gemini/Claude:** Yes; fix needed.
  - **GPT:** Agrees.
- **Status: Resolved: actionable fix.**
- **Evidence:**
  - "People claim" vs "real that happens" asymmetry.

H9. The-gap — Directional relief bias

- **Dispute:** Is the book an on-ramp from Protestant to high-density traditions?
- **Positions:**
  - **Gemini:** Yes; asymmetry in relief.
  - **Editor:** Intentional editorial choice.
- **Status: Resolved as design choice, not blind spot.**
- **Evidence:**
  - Editor's explicit intention; translation emphasis directionality.

H10. The-gap — Ecumenical consensus baseline

- **Dispute:** Shared prior or deliberate editorial stance?
- **Positions:**
  - **Models (Wave 3/4):** Shared prior; potential bias.
  - **Editor (Wave 4 notes):** Intentional choice.
- **Status: Resolved as design choice, not undetected bias.**
- **Evidence:**
  - Editor explicit; still contested in traditionalist circles.

H11. The-gap — Relief-framing stability problem

- **Dispute:** Does "translation resolves disagreement" flatten genuine doctrinal divides?
- **Positions:**
  - **Claude:** Yes; needs explicit scope limits.
  - **Gemini/GPT:** Accepts; recommends explicit boundaries.
- **Status: Open; needs manuscript-level clarification.**
- **Evidence:**
  - Confessional theologians would still disagree on substance.

I. Methodology Design Disputes

I1. Blind Stage 1 vs Pre-mortem requirement

- **Dispute:** Should editor's thesis be hidden to preserve independence?
- **Positions:**
  - **Gemini:** Blind Stage 1.
  - **Claude:** Pre-mortem; blind is impractical.
- **Status: Resolved as hybrid.**

- **Evidence:**
  - Blind is coordination-chaotic; pre-mortem forces adversarial case while keeping scope.

I2. Model diversity vs prompt diversity

- **Dispute:** Is multi-model value due to different models or just repeated prompts?
- **Positions:**
  - **Claude:** Empirical question; could be prompt diversity.
  - **Others:** Agree it's open.
- **Status: Open.**
- **Evidence:**
  - No within-project experiment isolating model vs prompt effects.

I3. Optimal model count

- **Dispute:** Is three optimal?
- **Positions:**
  - **All:** Unknown; diminishing returns likely.
- **Status: Open.**
- **Evidence:**
  - No controlled comparison.

I4. Minimum cross-review

- **Dispute:** How many cross-review passes are required?
- **Positions:**
  - **GPT:** One pass minimum; diminishing returns after.
  - **Others:** Converged.
- **Status: Resolved as one-pass minimum.**
- **Evidence:**
  - the-gap skipping cross-review correlated with more misses.

J. Meta-Convergence and Self-Assessment Risks

J1. Meta-convergence on pessimism

- **Dispute:** Was Wave 3's pessimism accurate or itself correlated?
- **Positions:**
  - **Wave 3 synthesis:** Cannot resolve; could be RLHF bias or true.
  - **Editor (Wave 4 notes):** Believes pessimism was overstated.
  - **Wave 5 synthesis:** Pessimism partly accurate, partly RLHF-driven; over-negative framing.
- **Status: Partially resolved; tension preserved.**
- **Evidence:**
  - Correlation evidence is real; framing skewed negative.

J2. "Method solves problems it creates"

- **Dispute:** Are many catches just fixing model-introduced errors?
- **Positions:**
  - **Claude hostile critique:** Yes; method introduces errors it later catches.
  - **Counter-argument:** Catching is still valuable; single model could also err without detection.
- **Status: Open.**
- **Evidence:**
  - Errors introduced by models (SSS, HIPAA) were caught by others.

K. Editor vs Models Tension (Wave 5 explicit)

K1. Models' caution vs editor's enthusiasm

- **Dispute:** Which perspective should define the method's truth?
- **Positions:**
  - **Models:** Correlation ceilings, bounded value, strategic blindness.
  - **Editor:** Transformative speed and deliverability; factual safety net is the point.
- **Status: Held in tension; no resolution.**
- **Evidence:**
  - Documented limitations vs lived counterfactual impact.

K2. Negativity audit (systematic pessimism)

- **Dispute:** Were Waves 1–4 systematically pessimistic?
- **Positions:**
  - **Editor:** Yes; over-negative framing.
  - **Wave 5 synthesis:** Yes, but not purely RLHF; mixed with real evidence.
- **Status: Resolved as partial bias.**
- **Evidence:**
  - Models flagged their own pessimism and did not correct it.

L. Open Questions (Unresolved Ledger)

1. **How correlated are frontier model errors, empirically?** (multi-model measurement lacking)
2. **Can corpus-level blind spots be detected at all without external experts?**
3. **Optimal model count and diminishing-returns curve.**
4. **Does prompt diversity substitute for model diversity?**
5. **Does insight generation degrade on a different timeline than error prevention?**
6. **Does convergence eventually collapse the method's value, or does amplification persist?**
7. **How should availability-value be measured (time saved, cost avoided, cognitive load, quality retained)?**
8. **Can operational knowledge gaps be systematically fixed (e.g., 3-AM reconstruction realities)?**
9. **Can hallucination taxonomy be validated for multi-model pipelines (not just single-model)?**
10. **Can editor-induced convergence be mitigated beyond the current hybrid pre-mortem?**

M. Method-Level Criticisms (Retained)

- **Shared-prior convergence is undetectable inside the method; editor is the only defense.**
- **Cross-review is weak at scope policing and strategic depth; strong at tactical error catching.**
- **Content drift and qualifier erosion are chronic; require mechanical tooling.**
- **Process depends on editor bandwidth; fatigue is a latent systemic risk.**
- **Pipeline lacks inherent guarantees unless all layers render fully (truncation/omission issues).**
- **Method may converge on balanced, centrist conclusions due to shared training priors.**

N. Evidence Disagreements That Remain

N1. Specifics of hallucination correlation

- **Dispute:** Which hallucination types are most correlated in practice?
- **Status: Open.**
- **Evidence:**
  - The taxonomy is mechanistic; direct multi-model tests are absent.

N2. "80–90% of specialist quality" claim

- **Dispute:** Is this quantification valid or rhetorical?

- **Status: Open.**
- **Evidence:**
  - No formal scoring; appears as heuristic in Wave 5 synthesis.

Closing Note

This document preserves **every criticism, disagreement, and open question** across five waves and two editorial notes. It is not a verdict. It is the catalog of tensions that should remain visible to any reader of the Swiss Cheese Method paper.